

Identifying and Fixing Performance and Logic Issues



Chris B. Behrens

SOFTWARE ARCHITECT

@chrisbbehrens



Logging and Instrumentation with SSIS



Logging Data in SSIS

We're going to attach a log provider to the event hooks that SSIS fires

The events always fire, but without a delegate, they go nowhere and do nothing



Logging Providers in SSIS



Sql Server

Text File

XML File

Event Log

Sql Profiler

I don't recommend Event Log or Sql Server...

Because those are very hard to log ship

The Text Providers



Text file (CSV)



Xml file



Sql Profiler



When in doubt,
use text file

The Sql Trace Provider

The ability to *replay the log*

This allows you to re-execute
the output of the package
against a target db



Demo



Configure logging for our package

- With the text file provider
- Select some events

Execute our package

- Review the logs
- Talk about what it means



An Anti-pattern for Integration Services Performance



A Massive Data File



The file had dozens of fields



The size of the file was around a gigabyte (and growing)



We load the entire file and table into memory



Then join them together to perform INSERT, UPDATE and soft-DELETE operations



How Our Approach Changed

We began with a small set of sample data

Our approach didn't hold up under the full size load

So, we loaded the entire file into a temporary table, and resolved it inside

Big Table



Demo



Modify our ETL package to import the data to a temporary table

Create a view that reconciles the imported data

Created a stored proc that imports the new data

Execute our new package

Talk about a problem with this approach



Half-splitting



A Kill to A View



Our view transforms the data types



But if the source data doesn't conform to the target data type...



The view will fail with an exception



And you have to find the needle in the record haystack




A Tool for Your Toolbox



We've got a tool for
this kind of problem



And it can be applied
widely to problems in
technology and
elsewhere



Half-splitting allows us
to isolate the problem

Half-splitting a View

A view is just a query

We can limit the scope of the query with different syntaxes

And figure out which half of the records contains the problem



Half-splitting Queries

```
SELECT TOP 3 Row = ROW_NUMBER() OVER (ORDER BY Id), ID, Title, Description, Points FROM vw_TempActivities
```

```
SELECT TOP 3 Row = ROW_NUMBER() OVER (ORDER BY Id), ID, Title, Description, Points FROM vw_TempActivities  
ORDER BY ROW desc
```

```
WITH expr AS
```

```
(  
    select    Id,  
              Row = row_number() OVER ( order by ID)  
    from      vw_tempActivities  
)
```

```
SELECT  A.ID, A.Title, A.Description, A.Points  
FROM    expr E JOIN vw_tempActivities A on E.Id = A.ID  
WHERE   Row > 3 AND Row < 6
```



Fixing the Problem

Once you have
found the record...

For example, an
“N” in what should
be a numeric
field...

You can decide
how to fix it



```
SELECT * FROM TEMP_Activities WHERE Points = 'N'
```

Multiple Bad Records

If there's one problem...

There are probably more

This is not a problem




If it's in both halves (or eighths, or 1/2048), just pick one, either one will work



The Worst Thing You Can Do



Why This Is so Bad



Exceptions Don't Belong in SSIS



The Nature of Exceptions



Exceptions should be exceptional



Of course, errors are going to happen (often)



But practically, the only solution is to redesign your process to eliminate it

Summary



Getting good information out of our system With logging from SSIS events

- OnError
- OnInformation

Log Providers

- Only three providers as useful
 - Text Provider
 - Xml Provider
 - Sql Trace

Common anti-pattern

- Big Table
- How to get around it
 - Temp tables
 - Views
 - Stored procs

