

Testing and Verifying an SSIS Package



Chris B. Behrens

SOFTWARE ARCHITECT


@chrisbbehrens



What We Mean By “Testing”



Aristotle and Unit Testing



The Nichomachean Ethics

Both his son and his dad were named Nichomachus

The measure of a thing is how well it fulfills its function

A good knife is a knife that cuts well

This philosophy applied to code tells us how to test

Testing with Databases



This approach is more difficult to apply to databases

We shouldn't waste our time testing the product specification

We should rely on the manufacturer

If that doesn't sound like a good idea, we should find another vendor



Verifying Availability

This isn't about
whether the
database is up
and running

We're concerned
that the database
works *correctly*

Verifying
availability is
important – just
not what we're
talking about now



What We ARE Testing For




Stored procs are
easy(er) to
verify



Execute the
proc and then
measure the
result




We're going to
verify changes
with a
measurement



By verifying the
state of a table
when we're
done

A Database Test Harness



A tricky data layer



I would restore the
database before each
test...




And then tear it down
for the next test



Where Our Target DB Is Running





Don't sweat it if you're not too deep into Docker yet

We're going to use Docker to build up and tear down our test environment

With this, you could do it hundreds of times per day if you needed to

Sql Server is available as a Docker image

On both Windows and Linux



Demo



Take a look at running container database

The command I used to create it

Provision an entirely new database from scratch

- Using a modified version of that command
- Connect to it in SSMS



A Testing and Verification Architecture



Order of Operations

1. Developer makes the change to the package
2. Developer checks in the change
3. The check-in triggers a build
4. The build pulls the new package from version control
5. The build executes the docker run command and creates a new instance of Sql server
6. The build executes a script to construct the test target db
7. The build executes any necessary schema upgrades
8. The build executes the package against the db
9. The build verifies that the results match expectations



Order of Operations

The build executes the docker run command and creates a new instance of Sql server



Order of Operations

The build executes any necessary schema upgrades

<https://app.pluralsight.com/library/courses/deploying-databases-octopus>



Order of Operations

The build verifies that the results match expectations



Demo



Build up the simple verification within our package

- With a row count transformation
- A script step to validate it
- Execute our package

Break it by changing the data file



Demo



Take the assertion that we've just created

Build a test platform to run it on

Using our build process

Execute our build

Verify our package



Going Deeper with this Pattern


Our assertion is
trivial

The content could
be entirely wrong,
and just have the
right row count

But this is a
jumping off point
for deeper testing



How This Would Really Work



Building everything from scratch every time is slow

- We've only done this to keep things simpler

If we cooked all this into a Dockerfile and built an image off of it, it would be much, much faster

<https://bit.ly/2zcasjB>

This pattern is not Sql Server specific

Other databases have been in Docker a lot longer



Summary



Develop a unified testing methodology for our package

- Using Docker images
- A scripted assertion
- Jenkins builds

